



Optimistic Shared Memory Dependence Tracing


Yanyan Jiang¹, Du Li², Chang Xu¹, Xiaoxing Ma¹ and Jian Lu¹

¹Nanjing University

²Carnegie Mellon University

Understanding Non-determinism

- Concurrent programs are non-deterministic
 - transient behavior on specific interleaving/schedule¹

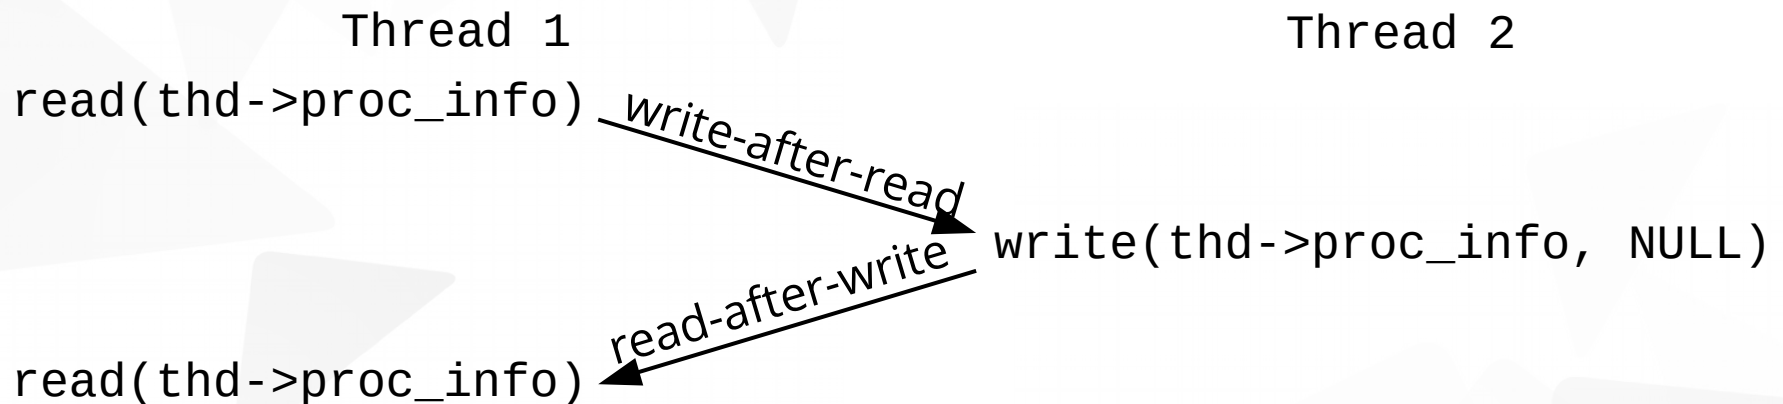


Understanding **order** of shared memory accesses

¹ MySQL crashes due to null pointer dereferencing.

Shared Memory Dependences

- ▼ The order between consecutive accesses of a shared location



- ▼ Four types of shared memory dependences
 - ▼ read-after-read (RAR), read-after-write (RAW)
 - ▼ write-after-read (WAR), write-after-write (WAW)

Using Shared Memory Dependences

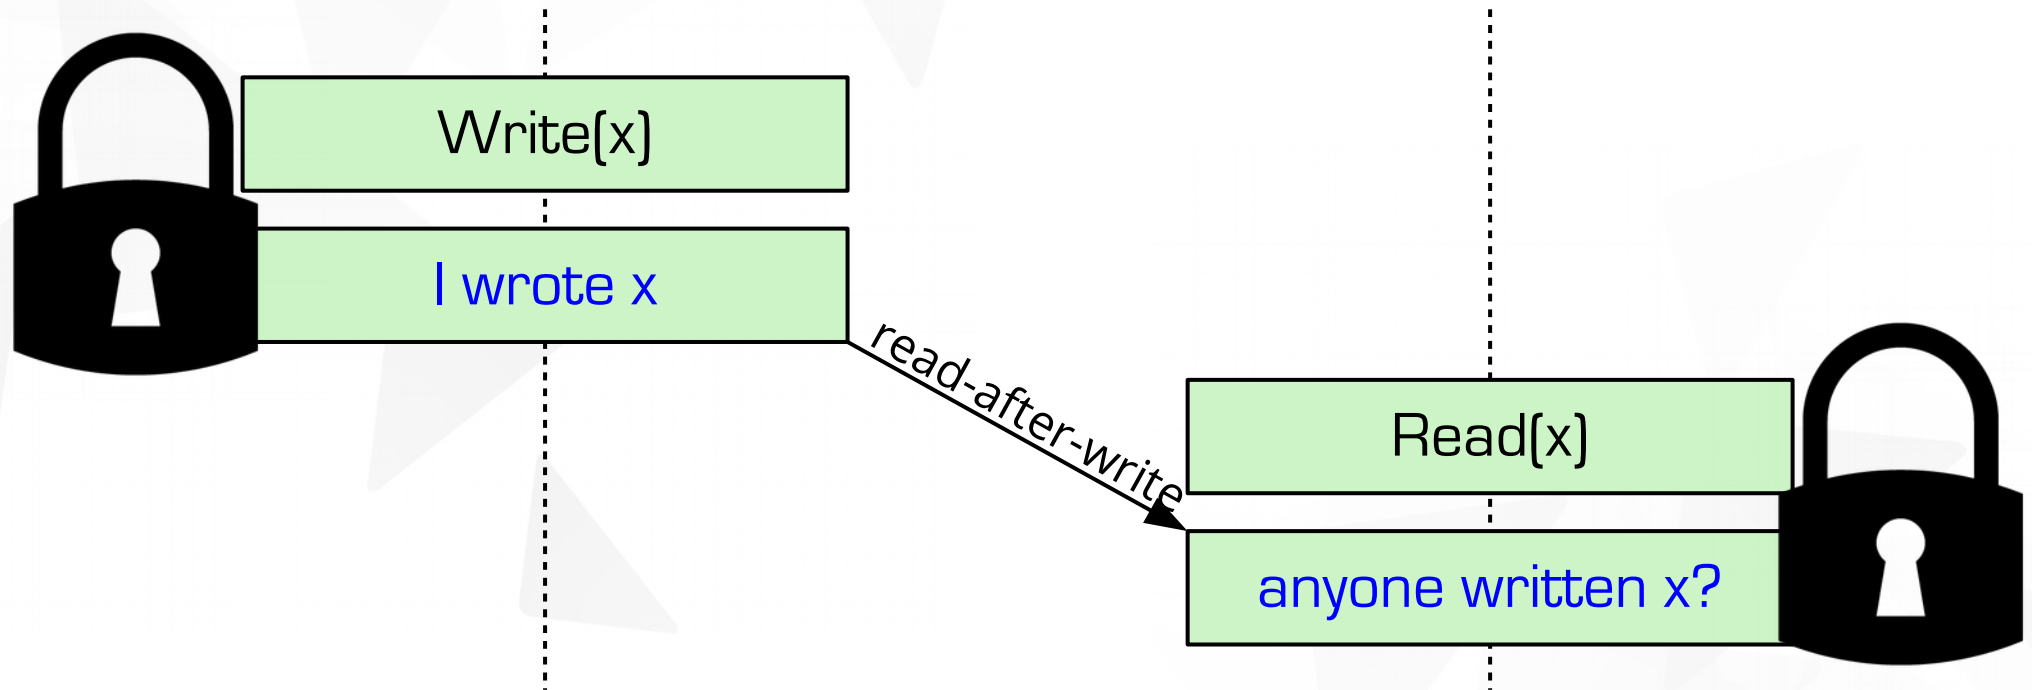
- ▶ Shared memory dependence in terms of replaying
 - ▶ RAW + WAR + WAW → trivial deterministic replay
 - ▶ RAW + WAW → $O(n)$ deterministic replay¹
 - ▶ RAW only → $O(\exp(n))$ deterministic replay²
- ▶ Predictive trace analyses
 - ▶ data race / atomicity violation detection

¹ Y Jiang, et al. CARE: Cache guided deterministic replay for concurrent Java programs. In *ICSE*, 2014.

² P Liu, et al. Light: Replay via tightly bounded recording. In *PLDI*, 2015.

Capturing Shared Memory Dependences

- Generally, we update metadata at shared memory accesses, and check it afterwards




We must force atomic access!

Capturing Shared Memory Dependences: Overhead

- ▼ Program instrumentation inevitably brings overhead
 - ▼ scalable overhead: thread-local operations
 - ▼ metadata bookkeeping
 - ▼ extra call and branch instructions
 - ▼ non-scalable overhead: serialization
 - ▼ forcing atomicity is much more costly

Reducing Overhead: Exploiting Thread Locality


- ▼ A thread tends to have exclusive/shared access of a shared location for a consecutive time period
- ▼ example: classical MESI cache coherence protocol



Fast path: exclusive/shared [CHEAP]
Slow path: bus traffic [EXPENSIVE]

Octet: Optimistic Tracing via Biased Locking¹

- ▼ Eager acquisition, lazy release



¹ M D Bond, et al. Octet: Capturing and controlling cross-thread dependences efficiently. In *OOPSLA*, 2013.

Octet: Drawback

- ▼ Biased lock is too optimistic
 - ▼ thread-local accesses are indeed fast, but
 - ▼ inter-thread coordination is too costly
 - ▼ under write-heavy workloads it is even worse than simple locking¹

Can we simultaneously achieve
FAST fast paths and **NOT-SLOW** slow paths?

¹ M Cao, et al. Drinking from both glasses: adaptively combining pessimistic and optimistic synchronization for efficient parallel runtime support. In *WoDet*, 2014.

RWTrace: Overview of the Trade-off

- Not only thread locality, but also *reads dominate writes*
 - slower thread-local writes
 - faster cross-thread coordinations

	LEAP ¹	Octet	RWTrace
Thread-local Read	Slow	Fast	Fast
Shared Read	Slow	Fast	Fast
Inter-thread Read	Slow	Very Slow	→ Slow
Thread-local Write	Slow	Fast	→ Slow
Inter-thread Write	Slow	Very Slow	→ Slow

¹ J Huang, et al. LEAP: Lightweight deterministic multi-processor replay of concurrent Java programs. In *FSE*, 2010.

1 – Serializing Writes

- ▼ Use simple mutex lock to protect writes
 - ▼ each address is associated with a lock
 - ▼ write-after-write dependences can be captured

	Octet	RWTrace
Thread-local Read	Fast	Fast
Shared Read	Fast	Fast
Inter-thread Read	Very Slow	Slow
Thread-local Write	Fast	Slow
Inter-thread Write	Very Slow	Slow

1 – Serializing Writes: Performance

- ▼ Write-write data race is NOT expected
 - ▼ often leads to unexpected behaviors
 - ▼ developers eliminate them by synchronization
- ▼ Corollary: write-time locking **scales** as the non-instrumented program scales!
 - ▼ serializing all writes by a global lock does not hurt performance too much¹

¹ J Zhou, et al. Stride: Search-based deterministic replay in polynomial time via bounded linkage. In *ICSE*, 2012.

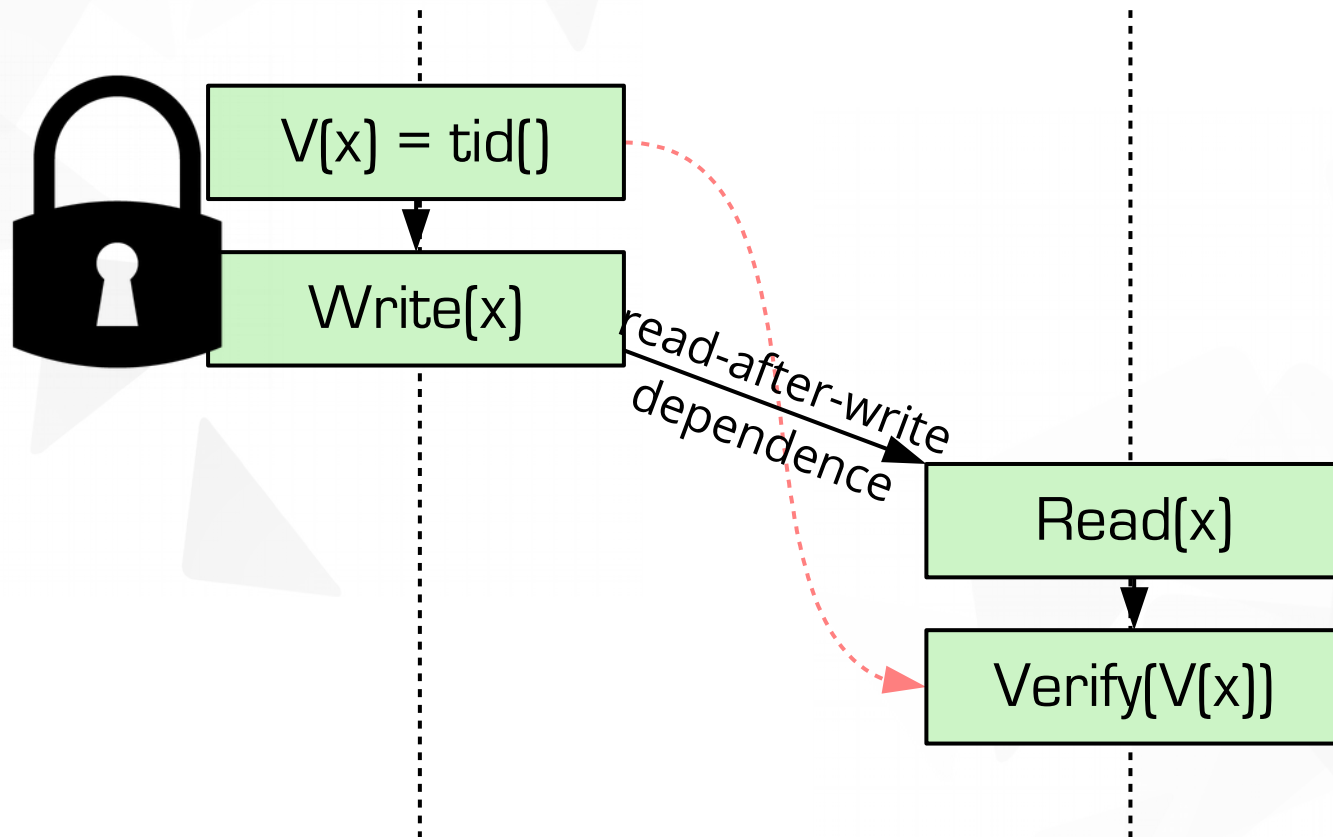
2 – Read Fast Path

- ▼ The vast majority, as fast as possible
 - ▼ $O(1)$, zero synchronization (wait-free), **scalable**
 - ▼ tests whether a thread is reading a value from previously unknown source

	Octet	RWTrace
Thread-local Read	Fast	Fast
Shared Read	Fast	Fast
Inter-thread Read	Very Slow	Slow
Thread-local Write	Fast	Slow
Inter-thread Write	Very Slow	Slow

2 – Read Fast Path: Am I Thread-local?

- Single-sided error: may unnecessarily fall back to slow path, but never miss any read-after-write dependence

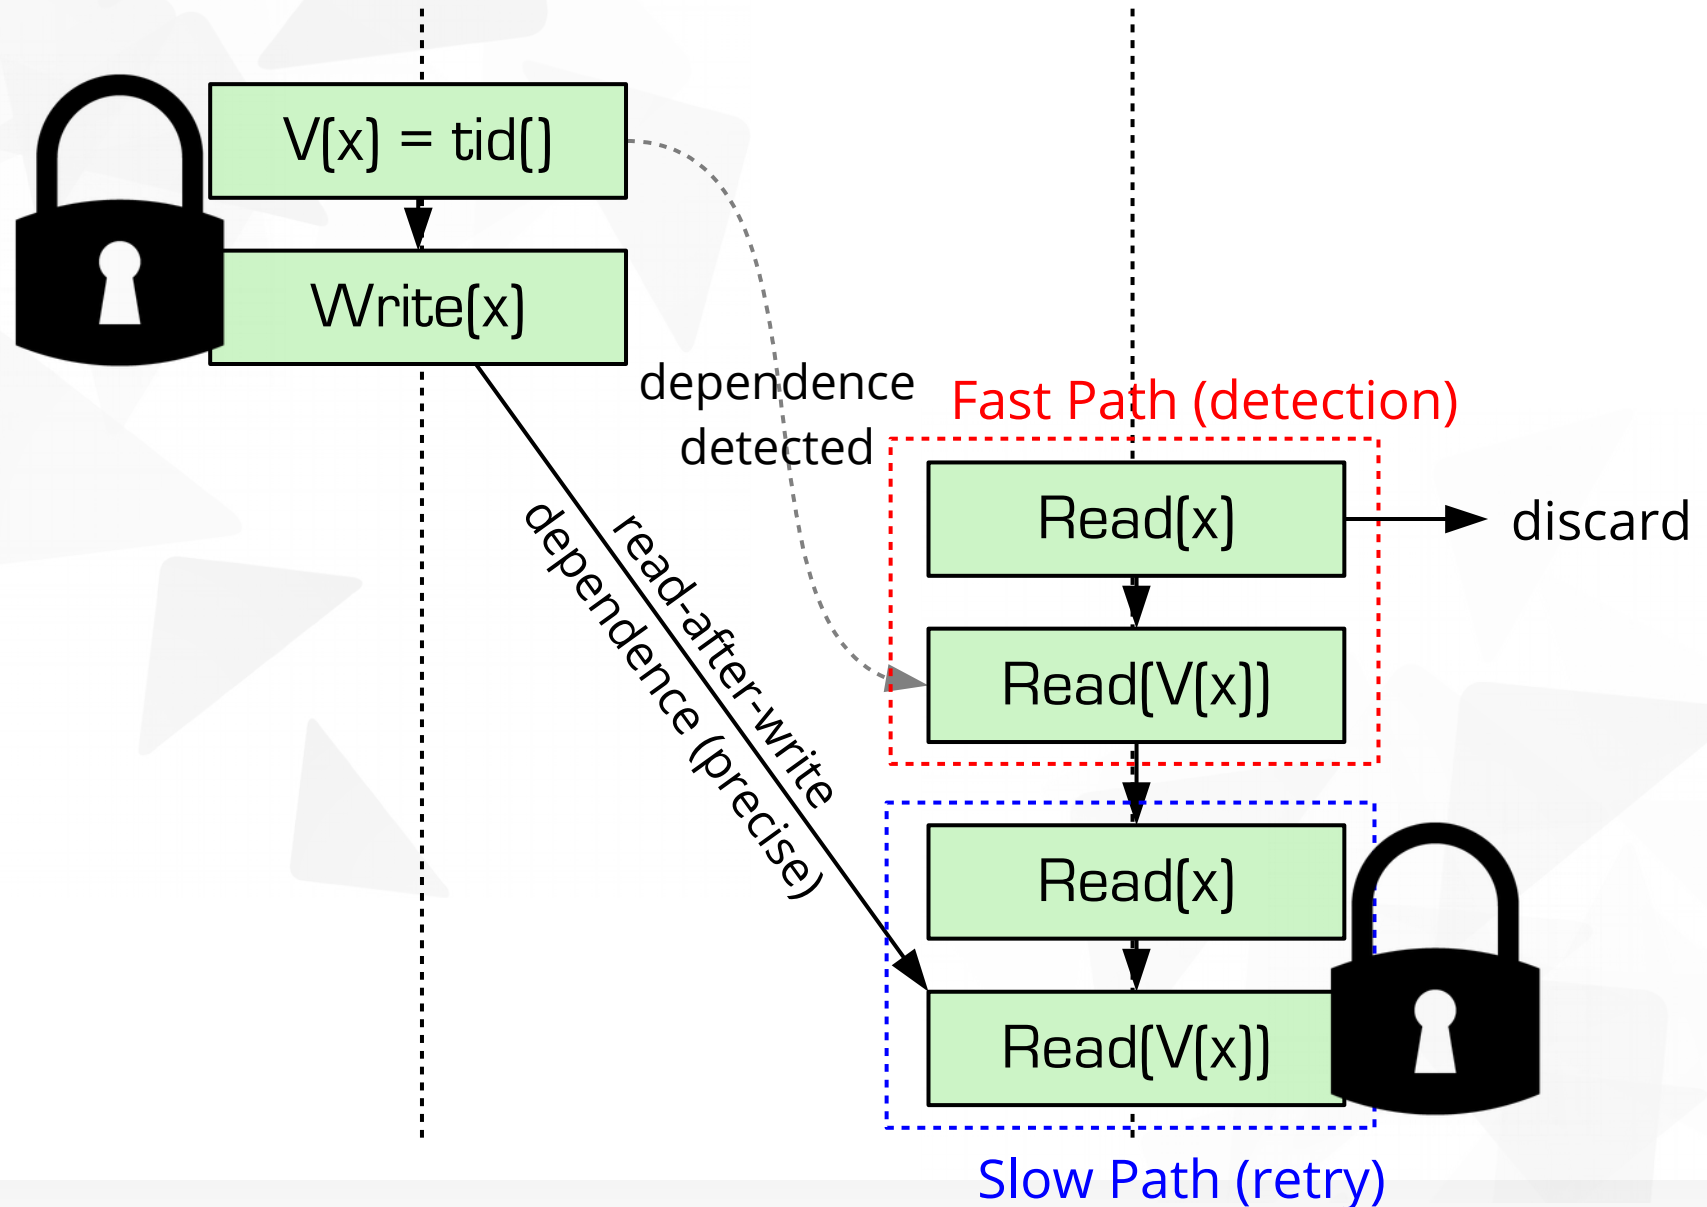


3 – Read Slow Path

- Read fast path fails → fall back to slow path
- ignore the previously read value
- read again with lock (exact read-after-write dependence)

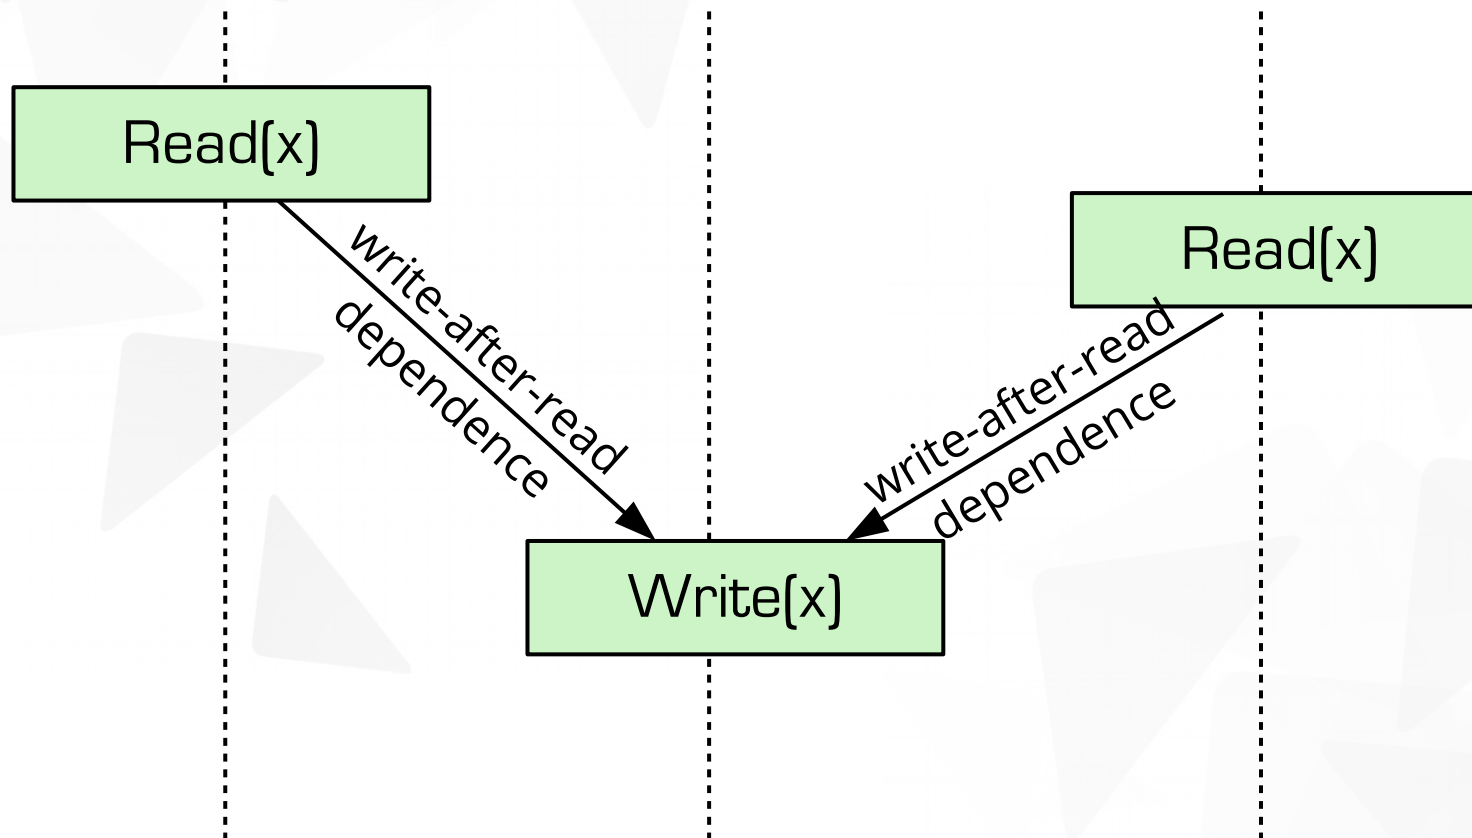
	Octet	RWTrace
Thread-local Read	Fast	Fast
Shared Read	Fast	Fast
Inter-thread Read	Very Slow → Slow	Slow
Thread-local Write	Fast	Slow
Inter-thread Write	Very Slow	Slow

3 – Read Slow Path: Retry



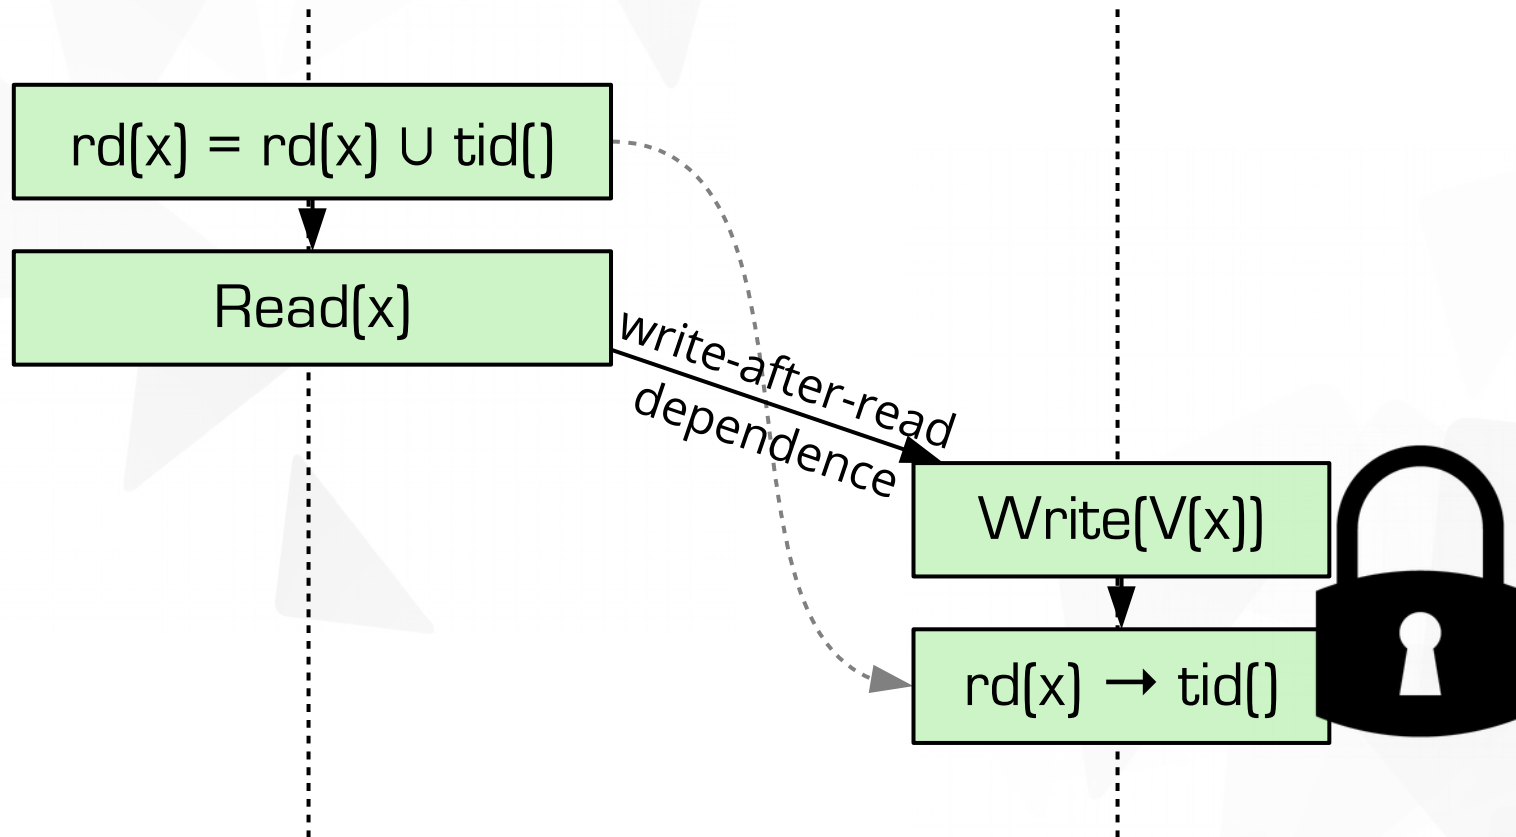
4 - Write-after-read Dependences

- Write-after-read dependences are many-to-one



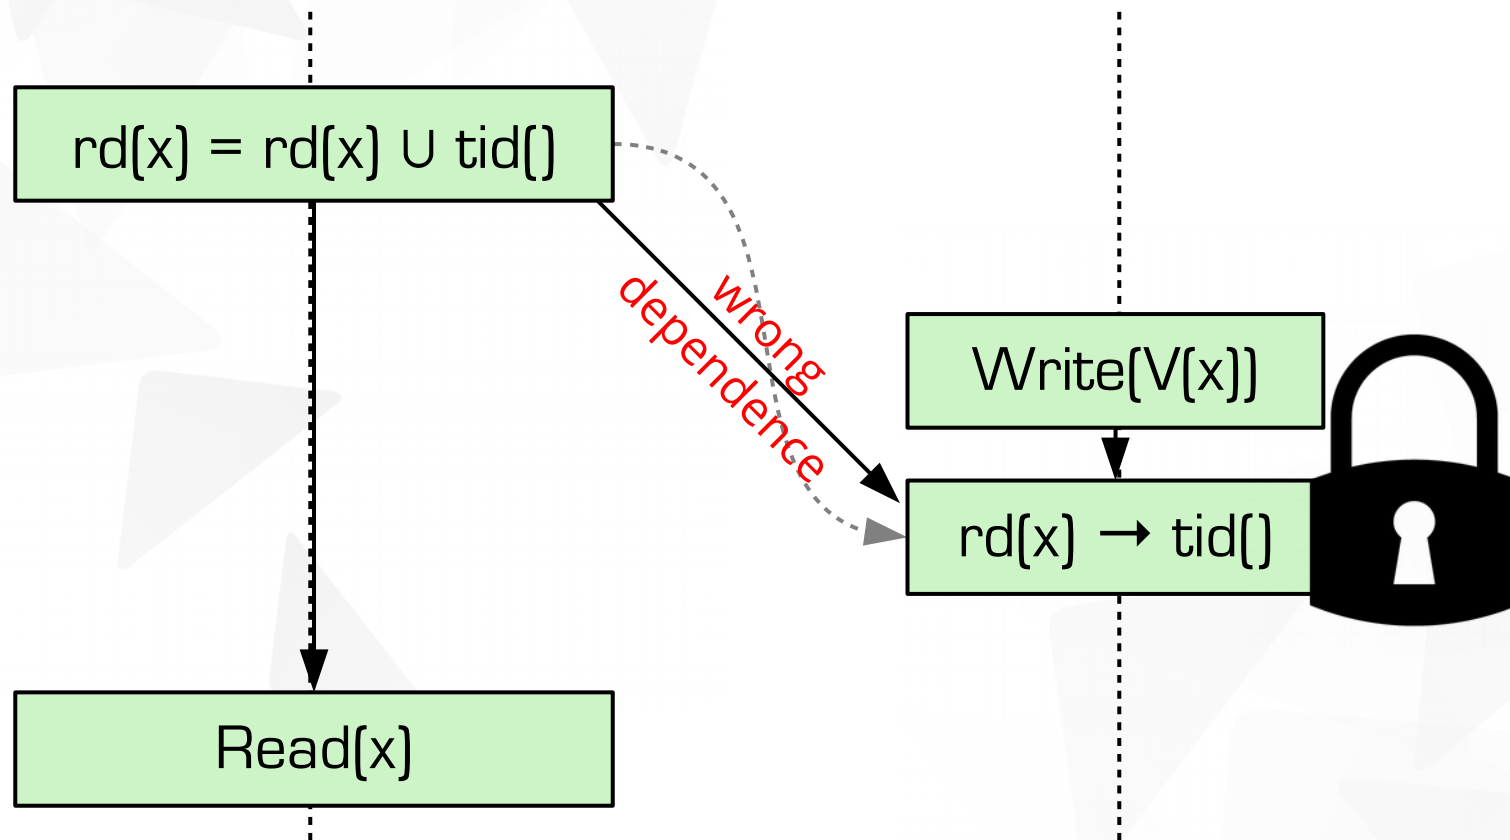
4 - Write-after-read Dependences: The Memory Order Trick Again, Inversely

- ▼ Maintain a read set to preserve wait-free read fast path
 - ▼ no extra synchronization, still **scalable**

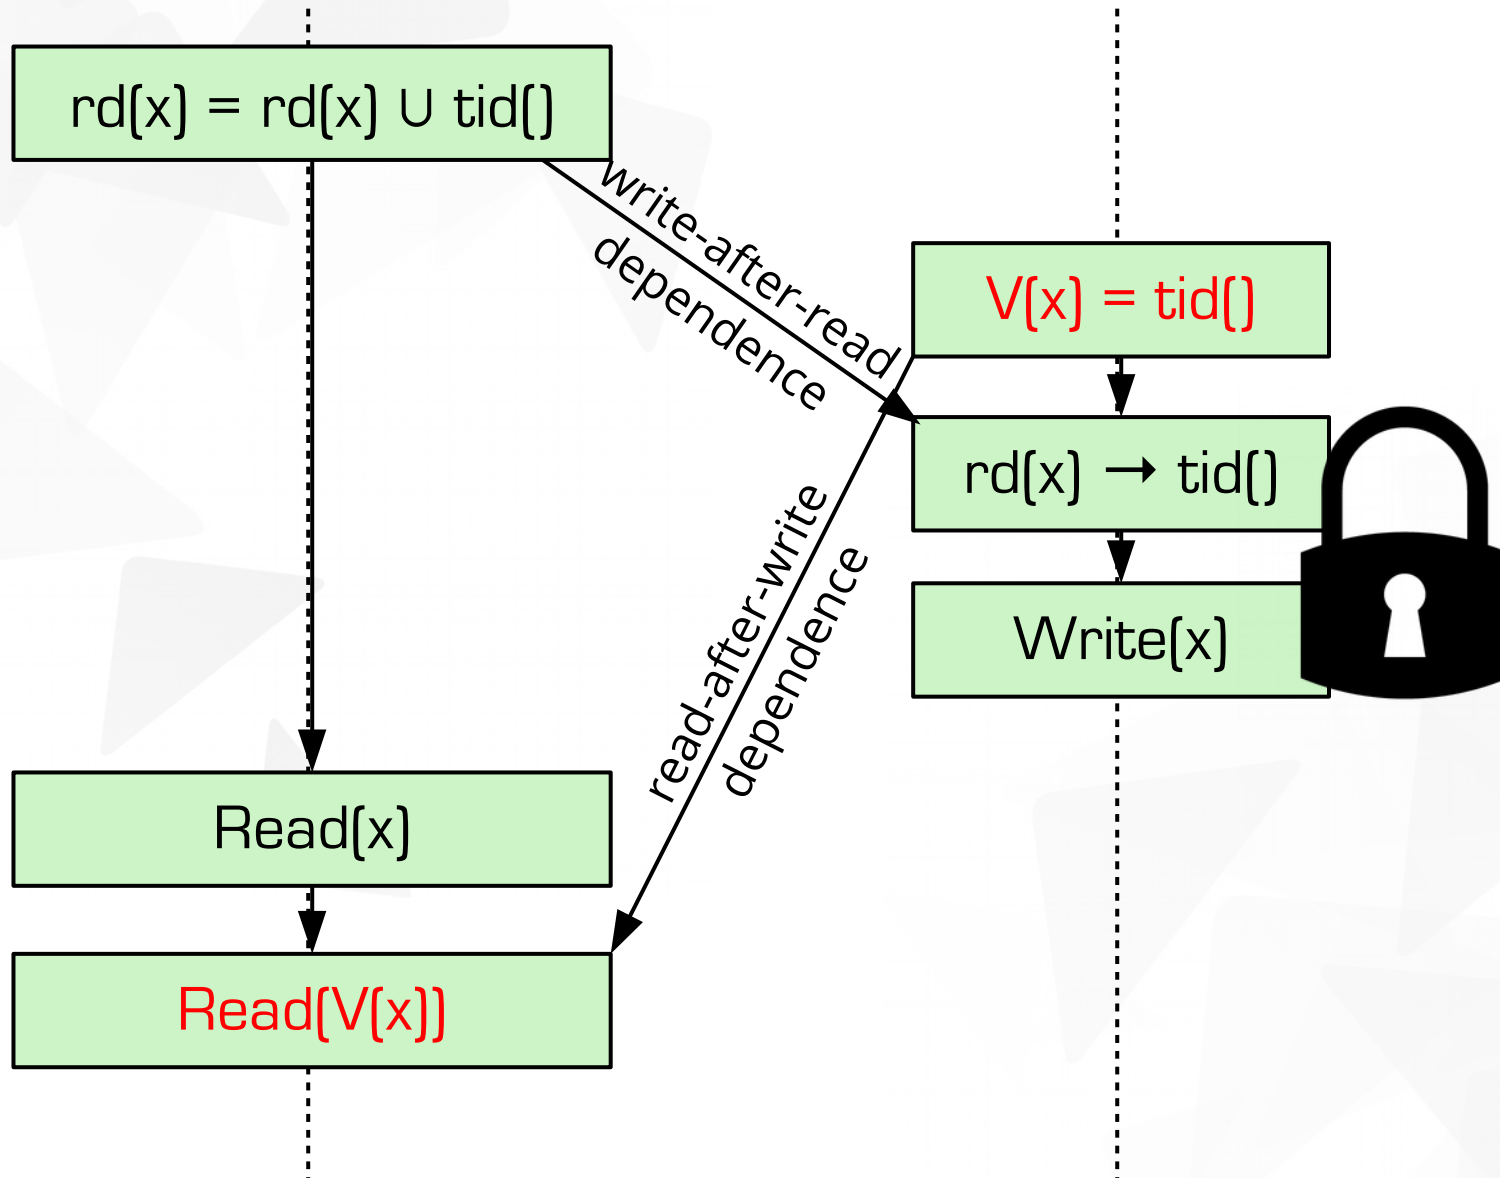


4 - Write-after-read Dependences: Wait-free Thread-local Read Fast Path

- There can be tricky memory orderings

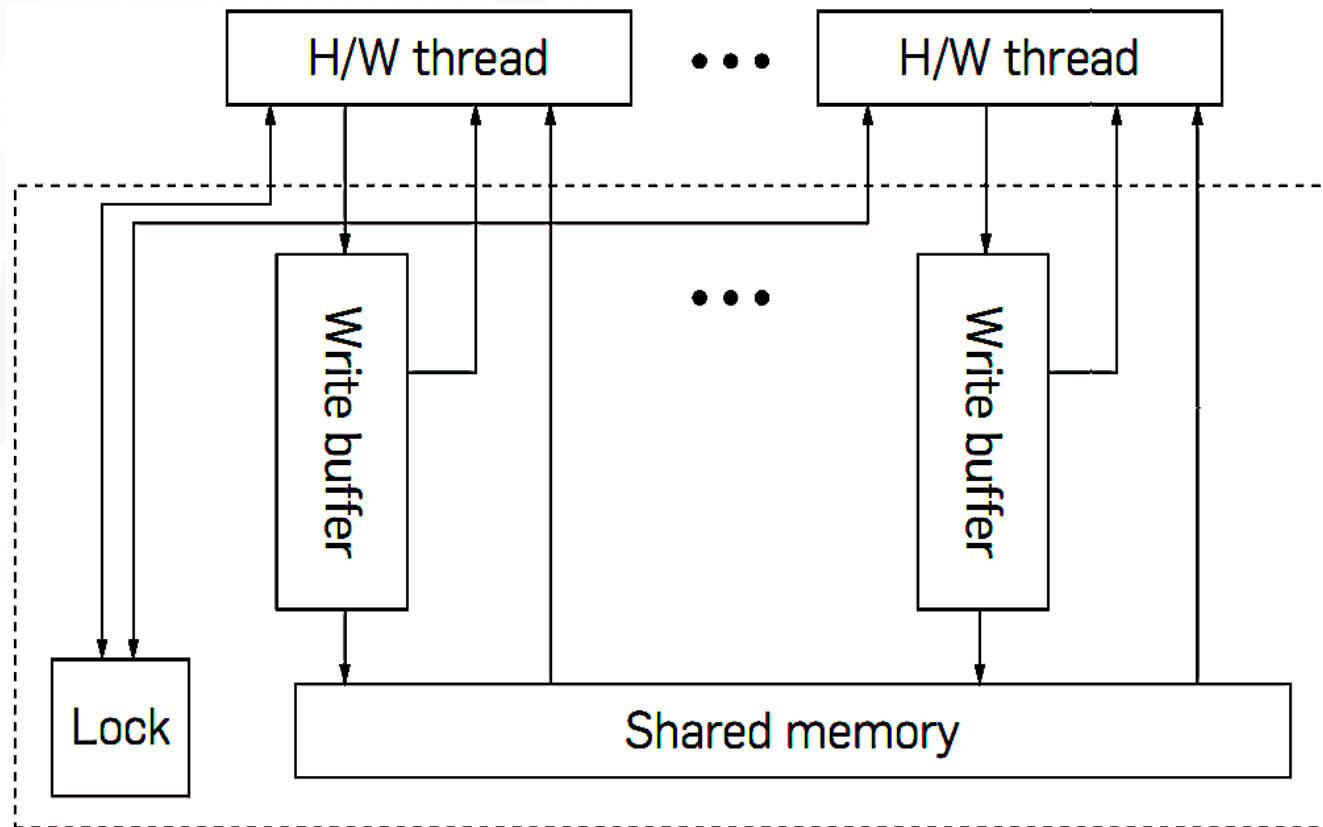


4 - Write-after-read Dependences: Handling Self-loop



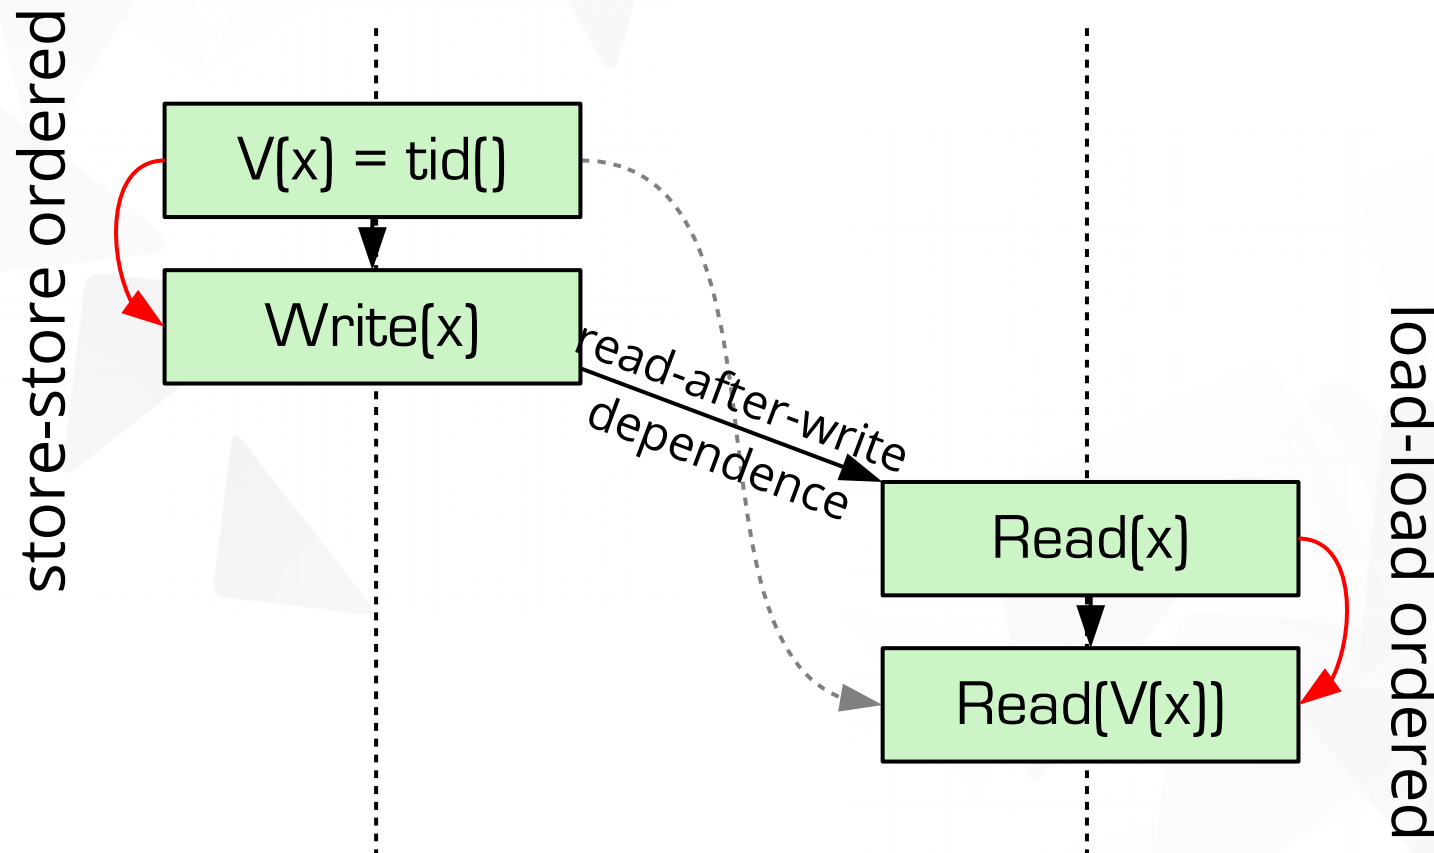
5 - Relaxed Memory Model

- ▶ RAW dependence tracing is barrier-free for x86-TSO
- ▶ fences required for weaker memory models



5 – Relaxed Memory Mode: Memory Ordering on x86-TSO

- Barrier-free for read-after-write dependences
- MFENCE for write-after-read dependences




RWTrace: Optimistic Shared Memory Dependence Tracing

- ▼ Technical highlights
 - ▼ precise WAW, RAW, WAR dependences tracing
 - ▼ wait-free thread-local read fast path
 - ▼ barrier-free on x86-TSO
 - ▼ scalable program instrumentation
- ▼ Implementation
 - ▼ upon LLVM, open source¹
 - ▼ as our experimental platform (like Octet)

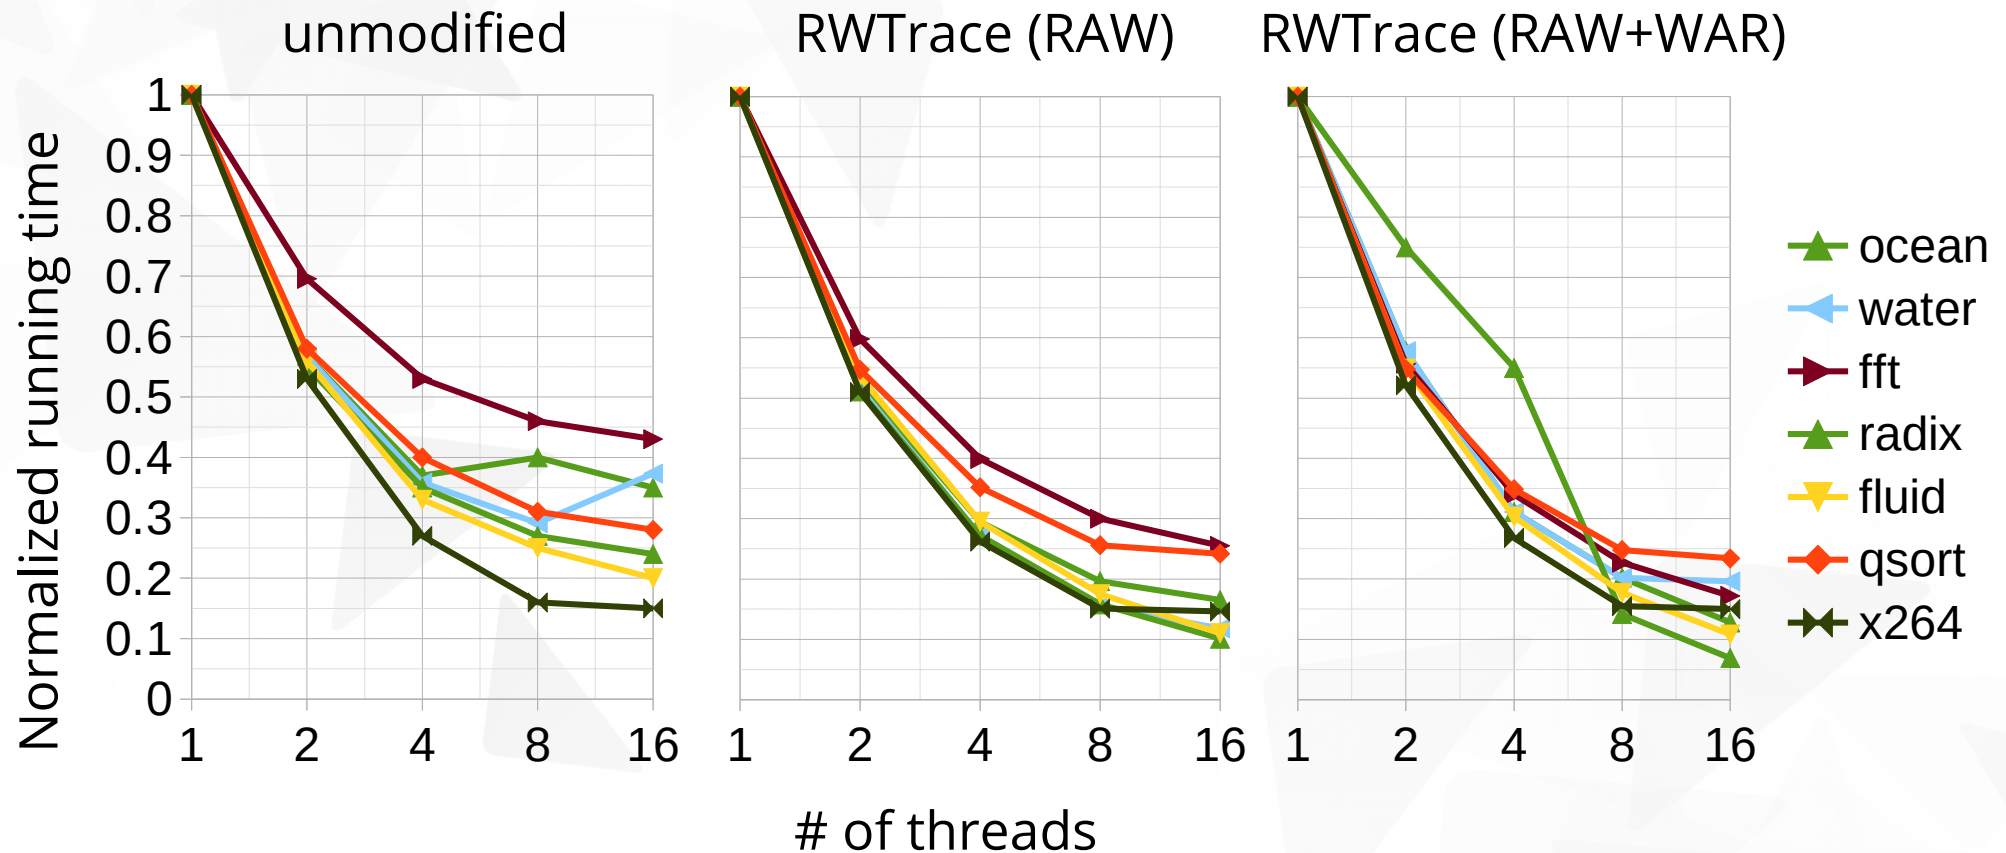
¹<http://github.com/jiangyy/rwtrace>

Evaluation Results¹: Overhead



¹ All experiments conducted on a 4×6-core Xeon machine.

Evaluation Results: Scalability



Shared Memory Dependences: In Pursuit of Determinism

- ▼ Dynamic analysis of concurrent systems
 - ▼ deterministic replay, data race / atomicity violation / false sharing detection, etc.
 - ▼ shared memory dependence reduction
 - ▼ theoretical aspects
- ▼ Non-determinism control
 - ▼ software transactional memory
 - ▼ deterministic multi-threading



Thank You!